

1 Softmax

Softmax is the function $\mathbb{R}^n \rightarrow \mathbb{R}^n$ given by

$$f_i(x) = \frac{e^{x_i}}{\sum_{k=1}^n e^{x_k}}$$

A bit of algebra shows that its derivatives are given by

$$\frac{\partial f_i}{\partial x_j} = \begin{cases} -\frac{e^{x_i} e^{x_j}}{(\sum_{k=1}^n e^{x_k})^2} & i \neq j \\ \left(\frac{e^{x_i}}{\sum_{k=1}^n e^{x_k}} \right) \left(1 - \frac{e^{x_i}}{\sum_{k=1}^n e^{x_k}} \right) & i = j \end{cases}$$

Note that this is symmetric. If we let $y = f(x)$ then this can be rewritten in vectorized form as

$$\frac{\partial f}{\partial x} = \text{diag}(y) - yy^T$$

Computing the product $(\frac{\partial f}{\partial x})z$ for a vector $z \in \mathbb{R}^n$ is given by

$$\left(\frac{\partial f}{\partial x} \right) z = y \circ z - yy^T z = y \circ (z - (y^T z) \mathbf{1})$$

where \circ is an elementwise product and $\mathbf{1} \in \mathbb{R}^n$ is a constant vector of ones.

Now suppose that $X \in \mathbb{R}^{n \times m}$ is a matrix of m inputs stored in columns, and $Y \in \mathbb{R}^{n \times m}$ is the matrix of outputs obtained by applying f to each column of X . Given a matrix $dY \in \mathbb{R}^{n \times m}$ of upstream derivatives, in the backprop step we must compute the matrix $dX \in \mathbb{R}^{n \times m}$ whose i th column is equal to $\frac{\partial f}{\partial x}$ evaluated at the i th column of X , multiplied by the i th column of Y . Using the above results, it is clear that

$$dX = Y \circ (dY - \mathbf{1} \mathbf{1}^T (Y \circ dY))$$

This can be efficiently implemented in numpy using broadcasting as:

$$dX = Y * (dY - \text{np.sum}(Y * dY, \text{axis}=0))$$